

Predictive Modeling of Student Academic Performance

Ansh Gautam

June 2026

Abstract

This formal report presents a comparative analysis of three supervised machine learning algorithms utilized for predicting student final grades (G3) based on historical performance and demographic data. Linear Regression, Decision Tree Regressor, and Random Forest Regressor models were trained and systematically evaluated using Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and the Coefficient of Determination (R^2). The findings demonstrate that linear assumptions hold significant predictive power in this educational context, with the Linear Regression model achieving the highest R^2 score of 0.8487. Ensemble methods also performed robustly, while individual decision trees struggled with variance and accuracy. This report outlines the methodology, theoretical underpinnings, detailed results, and discussions pertinent to implementing educational data mining solutions.

1. Introduction

In the contemporary educational landscape, the ability to forecast student academic outcomes is of paramount importance. Early identification of students at risk of underperforming allows educators and institutions to deploy targeted interventions, optimize resource allocation, and ultimately improve graduation rates. Educational Data

Mining (EDM) has emerged as a critical field, leveraging statistical and machine learning techniques to extract actionable insights from vast repositories of student data.

This project aims to build and evaluate predictive models capable of estimating a student's final grade (denoted as G3) based on a structured dataset containing prior academic records, study habits, and demographic variables. The central objective is to determine which algorithmic approach yields the highest accuracy and reliability for this specific continuous target variable.

To achieve this, three distinct machine learning paradigms were explored: parametric linear modeling (Linear Regression), non-parametric hierarchical partitioning (Decision Tree Regressor), and an ensemble learning technique (Random Forest Regressor). By comparing these models, this report elucidates the underlying data distribution characteristics and provides a clear recommendation for future predictive systems in academic settings.

2. Dataset and Preprocessing

The analysis relies on a comprehensive student performance dataset. Typical datasets in this educational domain include categorical variables (such as school support, family educational background, and extracurricular activities) alongside numerical variables (such as previous assessments G1 and G2, absences, and weekly study time).

Prior to model training, a rigorous preprocessing pipeline was implemented. This involved handling any missing values through mean or median imputation to preserve data integrity. Categorical variables were transformed into numerical formats using encoding techniques, ensuring compatibility with the chosen regression algorithms. Continuous variables, particularly those with wide ranges, were scaled to standardize the input space, a crucial step for distance-based and gradient-descent optimized models.

Finally, the dataset was partitioned into training and testing subsets, ensuring that the models are evaluated on unseen data. This process provides a realistic estimate of their generalization capabilities in a real-world educational environment.

3. Methodology

The predictive modeling phase involved the instantiation, training, and tuning of three distinct regressor models. Each model approaches the task of function approximation from a different mathematical perspective.

3.1 Linear Regression

Linear Regression serves as the baseline parametric model. It operates on the assumption that there is a linear relationship between the independent variables (features) and the dependent variable (target G3). The model seeks to find the optimal coefficients that minimize the residual sum of squares. The hypothesis function is defined as $Y = \beta_0 + \beta_1 X_1 + \dots + \beta_n X_n + \varepsilon$.

3.2 Decision Tree Regressor

The Decision Tree Regressor is a non-linear, non-parametric model that predicts the target variable by learning simple decision rules inferred from the data features. It recursively partitions the feature space into regions, assigning a constant predicted value to each leaf node. While highly interpretable, standard decision trees are highly susceptible to high variance and overfitting, especially on continuous data.

3.3 Random Forest Regressor

To mitigate the limitations of a single decision tree, the Random Forest Regressor employs an ensemble learning strategy based on bootstrap aggregating (bagging). It constructs a multitude of decision trees during training, each on a random subset of the data and features. The final prediction is the average of the predictions from all individual trees, an approach that significantly reduces model variance.

3.4 Evaluation Metrics

Model performance was quantified using three standard statistical metrics:

- **Mean Absolute Error (MAE):** Measures the average magnitude of the errors without considering their direction. Calculated as $MAE = (1/n) \sum |y_i - \hat{y}_i|$.
- **Root Mean Squared Error (RMSE):** A quadratic scoring rule that measures the average magnitude of the error. It gives a relatively high weight to large errors: $RMSE = \sqrt{(1/n) \sum (y_i - \hat{y}_i)^2}$.
- **Coefficient of Determination (R^2):** Represents the proportion of the variance in the dependent variable predictable from the features: $R^2 = 1 - (SS_{res} / SS_{tot})$.

4. Results

The models were evaluated on the held-out test set. The resulting performance metrics are summarized in the table below, highlighting the varying degrees of success each algorithm had in capturing the underlying patterns governing student performance.

Machine Learning Model	Mean Absolute Error (MAE)	Root Mean Squared Error (RMSE)	R^2 Score
Linear Regression	0.7651	1.2149	0.8487
Random Forest Regressor	0.7461	1.2409	0.8421
Decision Tree Regressor	0.9231	1.6502	0.7208

Key Observation: The Linear Regression model slightly outperformed the Random Forest ensemble in terms of variance explained (R^2), while the Random Forest achieved the tightest average error margin (lowest MAE).

The **Linear Regression** model demonstrated exceptional performance, achieving the highest R^2 score of 0.8487. This indicates that approximately 84.87% of the variance in the final grades can be explained by the model's linear combination of features. Its RMSE of 1.2149 was the lowest among the three models, suggesting fewer severe prediction errors.

The **Random Forest Regressor** provided highly competitive results, with an R^2 score of 0.8421, closely trailing the linear model. Notably, it achieved the lowest MAE of 0.7461, meaning its predictions were, on average, closest to the actual grades in absolute terms. This highlights the ensemble method's capability to balance bias and variance effectively.

Conversely, the **Decision Tree Regressor** exhibited the weakest performance profile. With an R^2 score of 0.7208, an MAE of 0.9231, and an RMSE of 1.6502, it was significantly outperformed by the other approaches. The sharp increase in RMSE compared to its MAE suggests that the model produced several predictions with large errors, a classic symptom of overfitting.

5. Discussion and Comparative Analysis

The empirical results provide compelling insights into the nature of the student performance dataset. The superior performance of the Linear Regression model strongly suggests a dominant linear relationship between certain features and the target variable, G3. In educational datasets, previous academic assessments (such as G1 and G2 scores) are notoriously highly correlated with the final assessment. Linear models excel in scenarios where the target is a direct, proportional continuation of historical numerical inputs.

While machine learning practitioners often default to complex, non-linear algorithms, this project underscores the principle of parsimony. The Random Forest model, despite its sophisticated aggregation mechanisms and ability to model complex interactions, could not decisively surpass the simple linear baseline in variance capture. Although it minimized the absolute average error (MAE), it could not capture the overall variance (R^2) quite as well as the linear model. This implies that the added complexity of the tree ensemble did not yield a commensurate gain in predictive power.

The pronounced underperformance of the standalone Decision Tree Regressor further illustrates the challenges of using single, highly flexible models on continuous educational data. Decision trees attempt to model continuous functions as step functions. When the underlying true function is strongly linear, a tree must create numerous, deeply nested splits to approximate it, inevitably capturing noise in the training set and failing to generalize well on the test set, as evidenced by its high RMSE.

6. Conclusion and Future Directions

This project successfully implemented and evaluated three predictive models for estimating student academic performance. The comparative analysis definitively shows that supervised machine learning can be a highly effective tool for educational forecasting. The Linear Regression model emerged as the most optimal solution for this specific dataset, offering the highest predictive accuracy ($R^2 = 0.8487$) coupled with the benefits of high interpretability and computational efficiency.

For future work, several avenues could be explored to further refine predictive capabilities. Feature engineering could be expanded; creating interaction terms or polynomial features might allow the linear model to capture subtle complexities. Deploying regularization techniques such as Ridge or Lasso regression could be investigated to confirm if further variance reduction is possible. Ultimately, integrating a robust model into a learning

management system could provide reliable, early-warning indicators to foster a more proactive educational environment.